

高性能计算程序设计 基础

任课教师：黄聃（Huang, Dan）

Based on the slides of Steve Leak

NERSC User Engagement Group

NERSC Cori


- 34 double-width cabinets
- 9,688 KNL + 2,388 Haswell nodes on Aries High-Speed Network
- 658,784 KNL cores + 76,416 Haswell cores
- Top500 #6 (June 2017)



Agenda

- Cori overview, logging in
- Run a simple job
- Building and running applications on Cori
 - Serial
 - Parallel (MPI)
 - Multithreaded (OpenMP)
- **What affects performance?**
 - Bottlenecks
 - Task placement and affinity
- **Preparing for performance analysis**

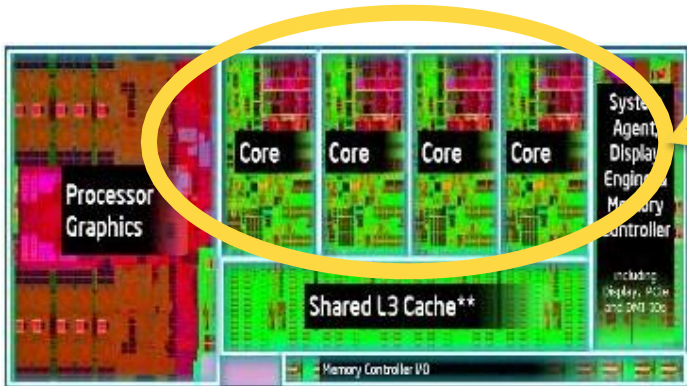
Cori Overview




- Login and compute nodes are distinct
- Large, fast, parallel
 - \$SCRATCH filesystem for running jobs
- Smaller \$HOME, configured for building code
- Burst buffer filesystem integrated, on high-speed network

What's so special about it?

- Your quad-core desktop CPU looks something like this:




- Compared to a Cori Haswell node:



16 cores
x 2 sockets,
128GB RAM

What's so special about it?

- And then there are the KNL nodes (68 cores, MCDRAM, ...)



What's so special about it?

- **But high-end CPUs don't make a supercomputer**
 - High speed interconnects between them
 - Lightweight compute node OS
 - Very large (28,000 TB) fast parallel filesystem
- **...and a different usage model**
 - Subset of nodes dedicated to a single task, run via batch system (no interactive GUI / desktop)

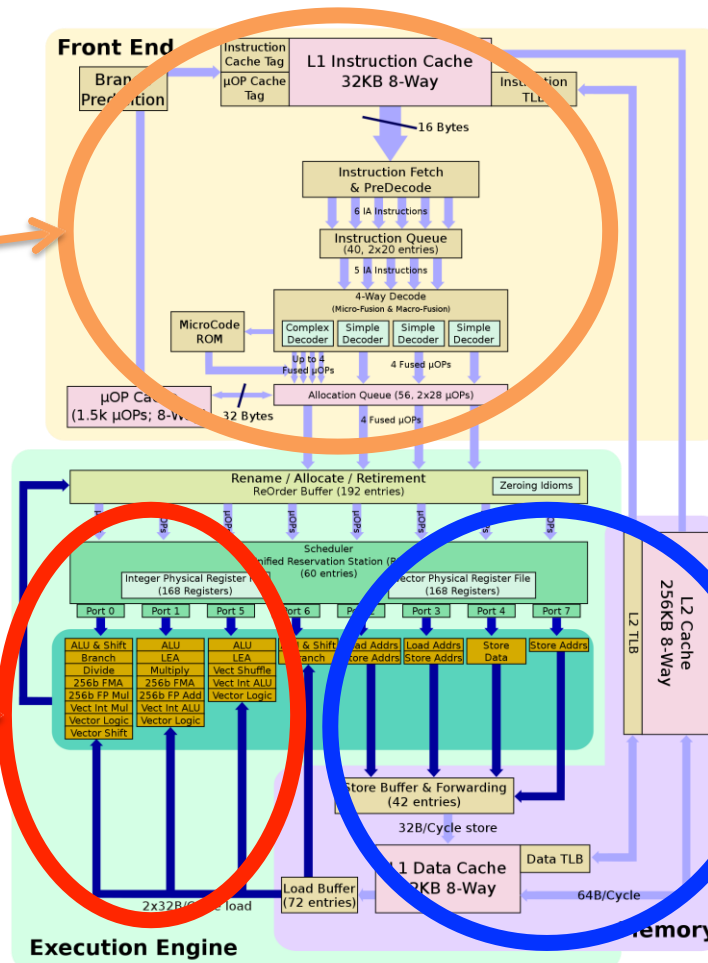
What's special about KNL?

- Different choice of compromise between die space allocated to different parts of the CPU

Instruction fetch-and-decode, hyperthreading, branching (important for eg compiling, GUI applications)

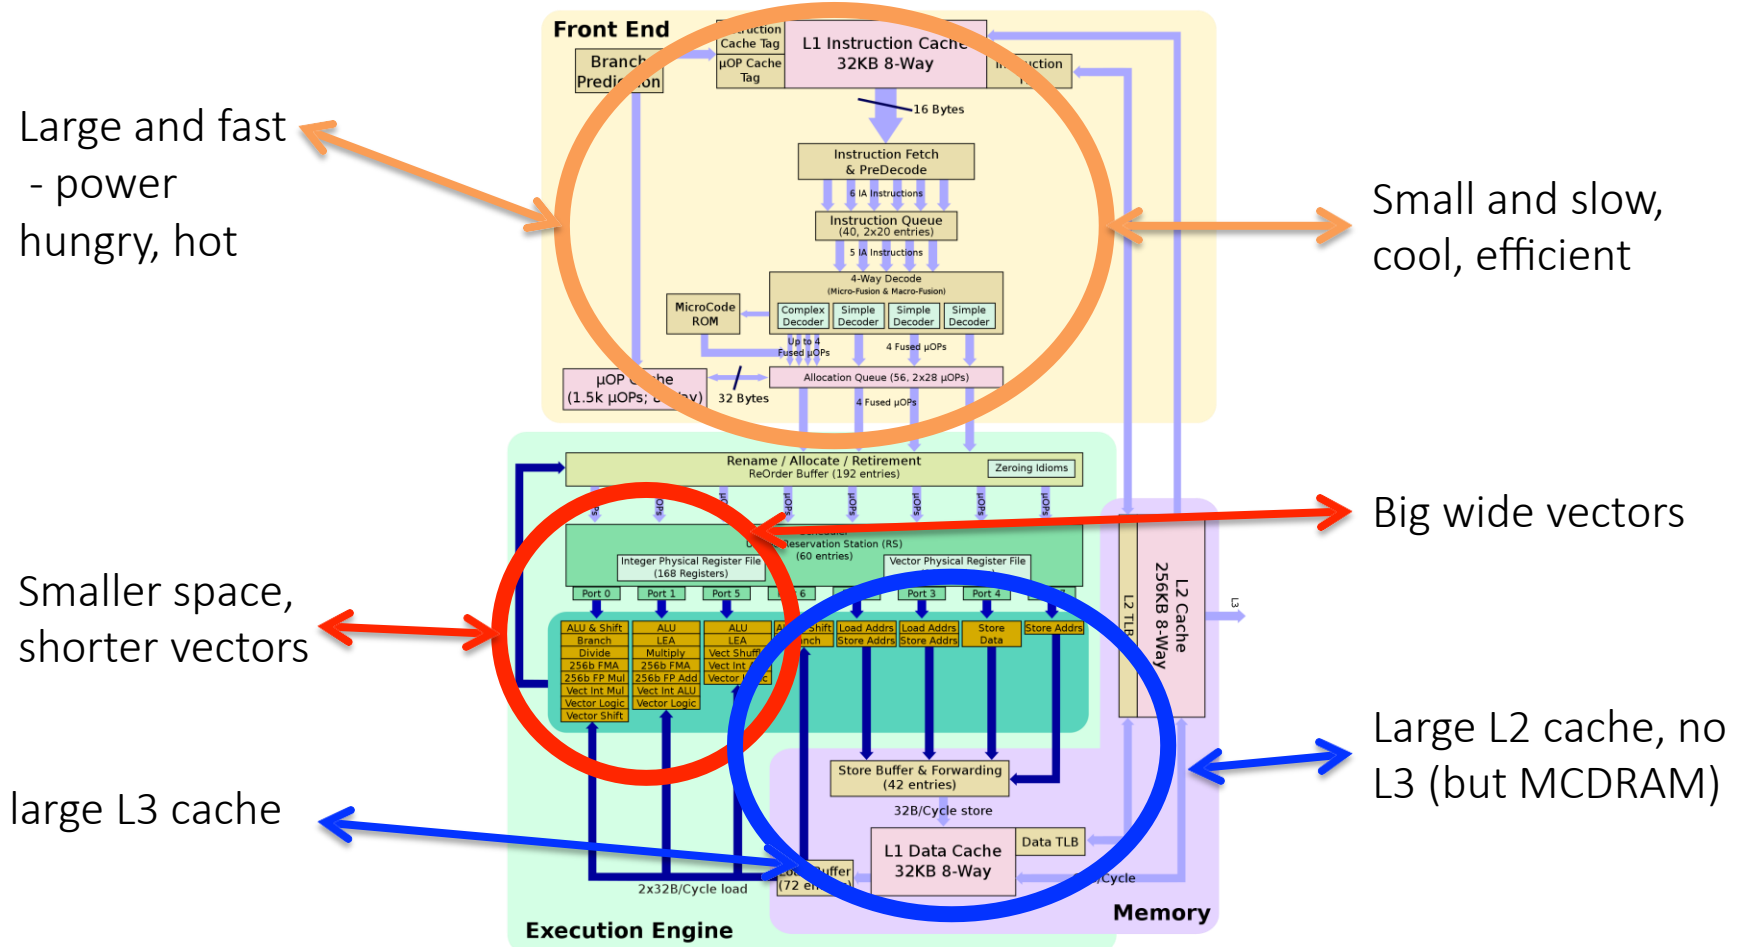
Arithmetic, vector and floating point units – the actual FLOPS

Memory access – keeps execution engine busy



What's special about KNL?


- Xeon (eg Haswell)
- Xeon Phi (KNL)



Why?

- **Exascale challenges – power and heat**
 - CPU frequency plateaued ~15 years ago
 - Transistor density, feature size reaching fundamental limits
 - Power consumption and heat dissipation are now the key constraints for supercomputing
 - ... we can't get there from here!
- **KNL emphasizes vectorization and parallelism at lower power**
 - Targets scientific computing

Cori Overview



- Login and compute nodes are distinct
- Large, fast, parallel
 - \$SCRATCH filesystem for running jobs
- Smaller \$HOME, configured for building code
- Burst buffer filesystem integrated, on high-speed network

Connecting to Cori

- If you have a UNIX-like computer, you can directly contact NERSC with your built-in SSH client
 1. Open a new terminal
 2. % `ssh -Y -<training_acct_username> cori.nersc.gov`
-Y Enables trusted X11 forwarding.
- Many SSH clients exist for Windows
 - A very popular one is **putty**
 - <http://www.putty.org/>
 - Advanced users might prefer to use SSH directly within **mintty** (from Cygwin distribution)

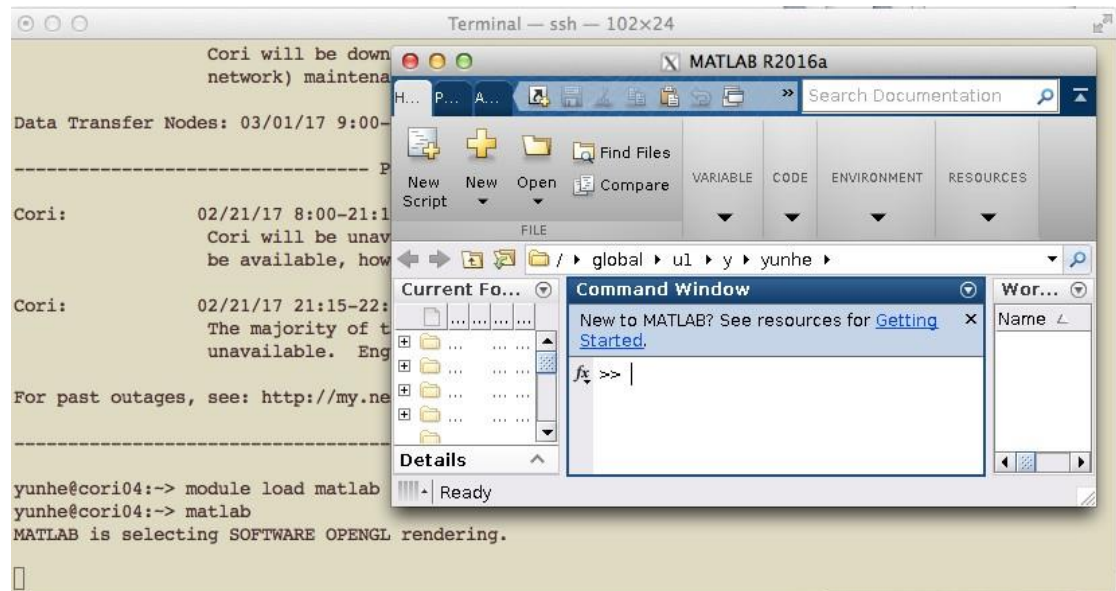
X-forwarding

- Allows you to access GUI programs remotely

Example:

```
localhost% ssh -Y -l elvis cori.nersc.gov
```

```
...  
e/elvis> module load matlab e/elvis>  
matlab  
<MATLAB starts up>
```



Example Session

```
localhost:~elvis> ssh -Y -l <training_account_name> cori.nersc.gov
```

```
*****
*
*                               *
*      NOTICE TO USERS      *
*      -----               *
*
*  Lawrence Berkeley National Laboratory operates this
*  computer system under contract to the U.S. Department of
*  Energy.      This computer system is the property of the United
*  States Government and is for authorized use only.      *Users
*  (authorized or unauthorized) have no explicit or implicit
*  expectation of privacy.*
*
*  Any or all uses of this system and all files on this system may be intercepted,
*  monitored, recorded, copied, audited, inspected, and disclosed to site, Department of
*  Energy, and law enforcement personnel, as well as authorized officials of other
*  agencies, both domestic and foreign.      *By using this
*  system, the user consents to such interception, monitoring, recording, copying,
*  auditing, inspection, and disclosure at the discretion of authorized site or Department
*  of Energy personnel.*
*
*
*  *Unauthorized or improper use of this system may result in administrative disciplinary
*  action and civil and criminal penalties. _By continuing to use this system you indicate
*  your awareness of and consent to these terms and conditions of use.      LOG OFF
*  IMMEDIATELY if you do not agree to the conditions stated in this warning._*
*
*
*****
```


Prompt on local system

Notification of acceptable use.

Password: <enter your training account password here>

Password prompt

After logging in...



- On a login node
 - cori01, cori02, ...
 - Shared by many users
 - Not necessarily the same one each time!
 - But same access to filesystems
- No direct access to compute nodes
 - Only via batch system (salloc, sbatch)
- Haswell (Xeon) architecture

Agenda

- Cori overview, logging in
- **Run a simple job**
- Building and running applications on Cori
 - Serial
 - Parallel (MPI)
 - Multithreaded (OpenMP)
- What affects performance?
 - Bottlenecks
 - Task placement and affinity
- Preparing for performance analysis

Running jobs – key points

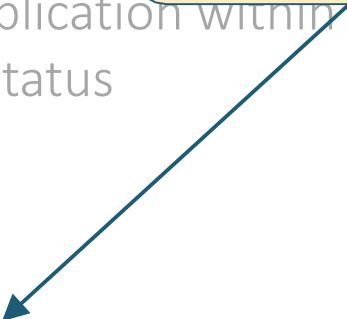
- **HPC work is via batch system**
 - Dedicated subset of compute resources
 - Login nodes are shared resource for building code, editing scripts, etc. Use batch jobs for real work
- **Key commands:**
 - sbatch / salloc - submit a job
 - srun - start an (optionally MPI) application within a job
 - sqs - check the queue for my job status
- **For today, we have a reservation**
 - #SBATCH –reservation=csgrrain

<https://docs.nersc.gov/jobs/>

Running jobs – key points

- HPC work is via batch system
 - Dedicated subset of compute resources
 - Login nodes are shared resource for building code, editing scripts, etc. Use batch jobs for real work
- **Key commands:**
 - sbatch / salloc - submit a job
 - srun - start an (optionally MPI) application within a job
 - sqs - check the queue for my job status
- **For today, we have a reservation**
 - #SBATCH –reservation=csgrrain

All of this is
on the web!




<https://docs.nersc.gov/jobs/>

How jobs work


Desktop / login node

- Timeslicing
 - core shared by multiple tasks
 - Works when the computer is mostly waiting for you



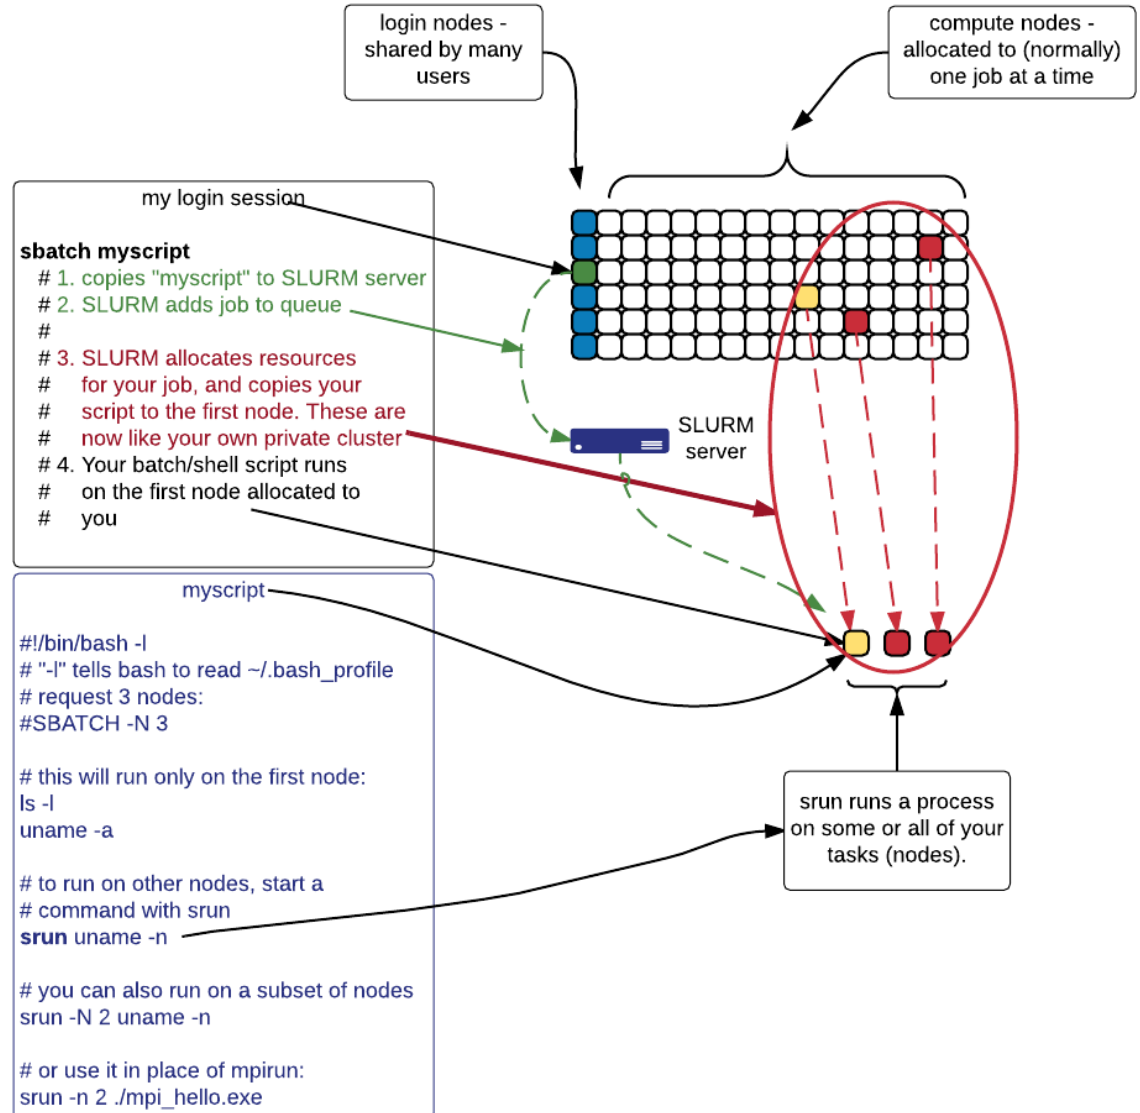
HPC

- You are waiting for the computer
- Subset of pooled resources dedicated to one job



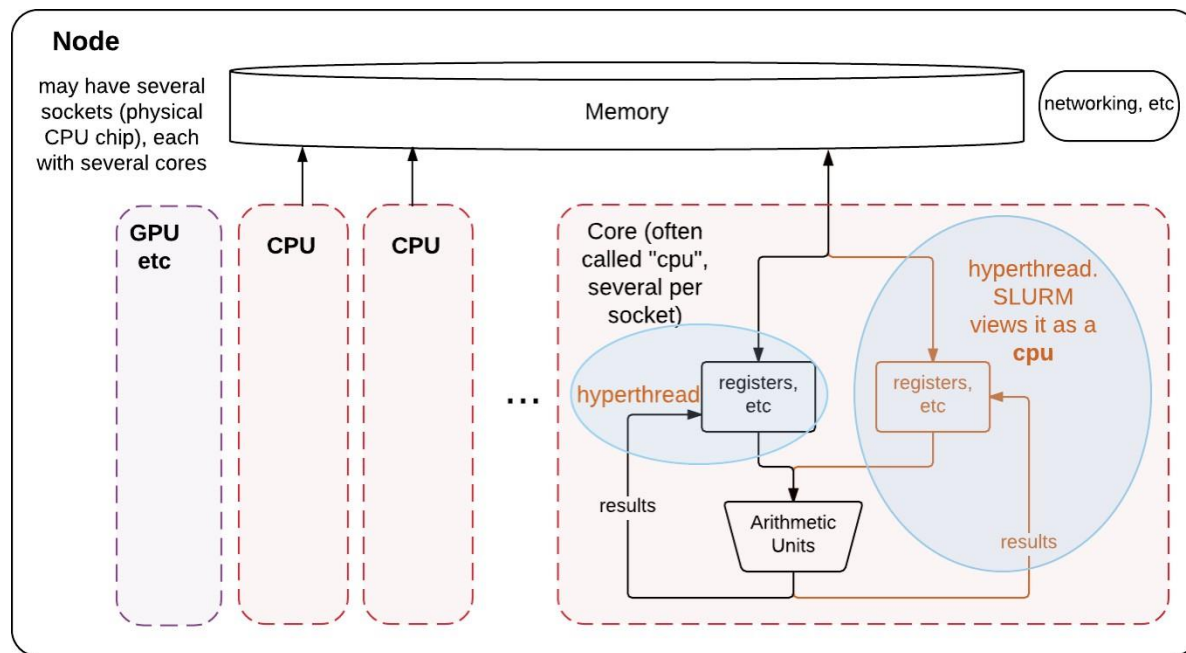
How jobs work

- Start on login node
 - shared by many users,
not for computational work
- Access compute nodes with sbatch or salloc
- Batch script
 - Copied to queue
 - Has directives for SLURM, and shell commands to perform on first compute node
- Access your other allocated nodes with srun
- stdout, stderr saved to file
 - (when running in batch mode)



Nodes, cores, CPUs, threads, tasks - some definitions


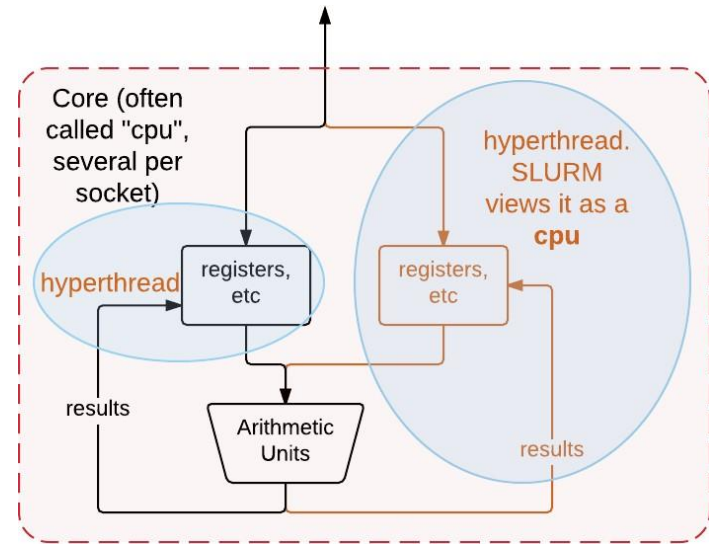
- *Node* is the basic unit of allocation at NERSC
 - Think “one host” or “one server”
 - Single memory space, multiple CPU cores (24 or 32 or 68 ...
 - And a core might support hyperthreading



Nodes, cores, CPUs, threads, tasks - some definitions


Hyperthreading

- Fast timeslicing
 - Good when arithmetic units frequently wait on memory
- Core holds state of 2 (4 on KNL) processes, they share arithmetic units
- **SLURM views each hyperthread as a CPU**
- But most HPC jobs perform best when not sharing a core!
- Usually best to reserve 2 (or 4) CPUs / core



context switch
(overhead)

Without
hyperthreading

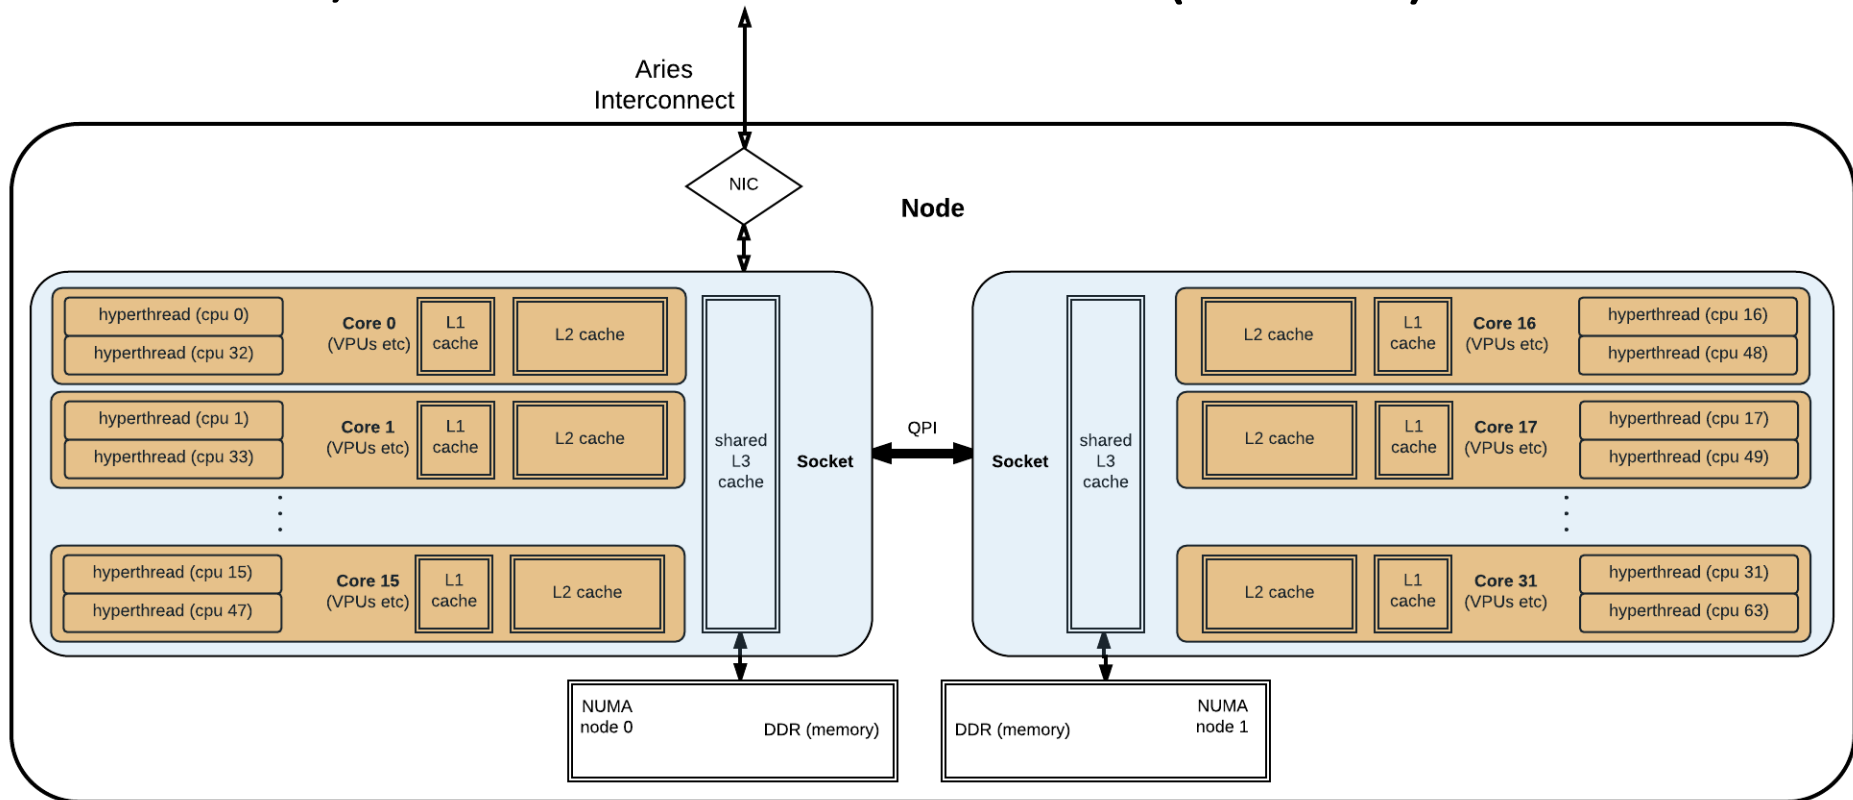


context switch
(overhead)

With
hyperthreading

Slurm tasks

- First, a block diagram of how hyperthreads, cores, cache, and sockets relate within a (haswell) node



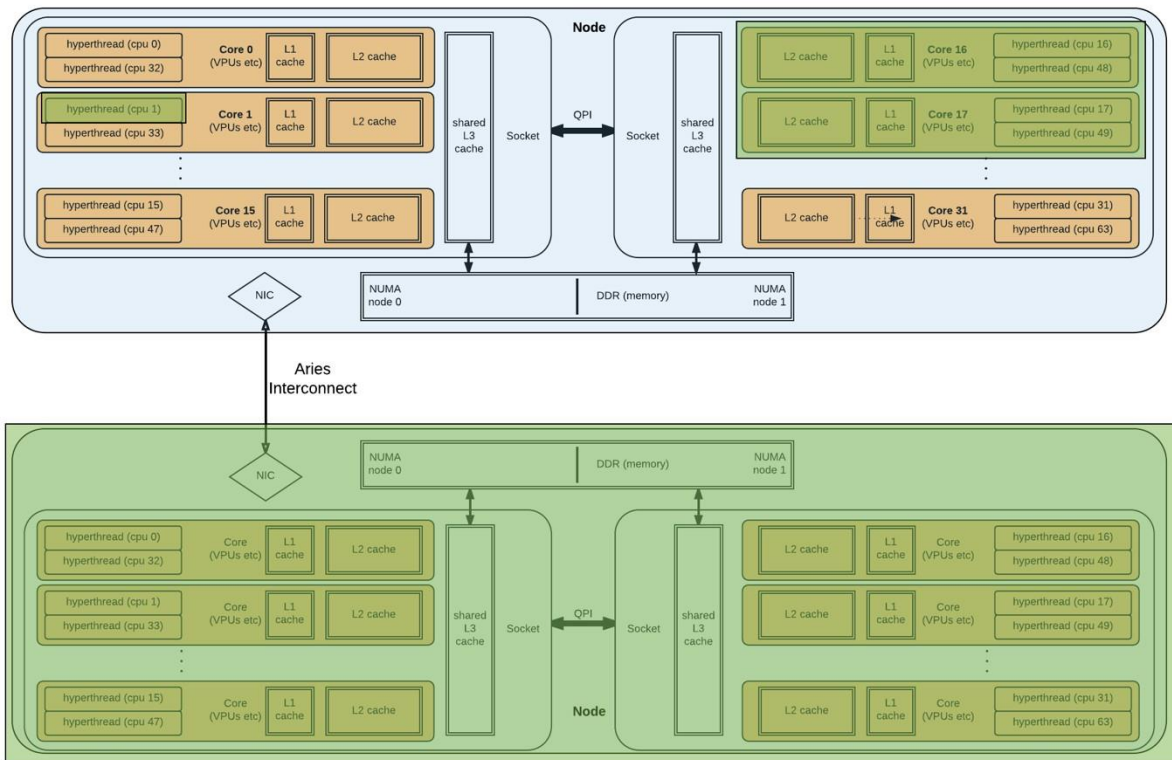
Slurm tasks

- A SLURM *task* is a reservation of CPUs and memory, up to one full node

- A *job* has many tasks
- 1 task typically corresponds to 1 MPI rank

`srun -n <ntasks>`

- Eg: 3 possible tasks on 2 nodes



So what must I request for my job?

What the batch system needs to know:

- How many nodes (or CPUs or tasks) does this job need?
- For how long does it need them?
 - Wallclock time limit

NERSC-specific extras:

- What type of CPU? (-C ...)
 - KNL or Xeon (haswell/ivybridge)?
- Which filesystems will this job use? (-L ...)
 - Usually SCRATCH

Requesting nodes or tasks

```
#SBATCH -N 64
```

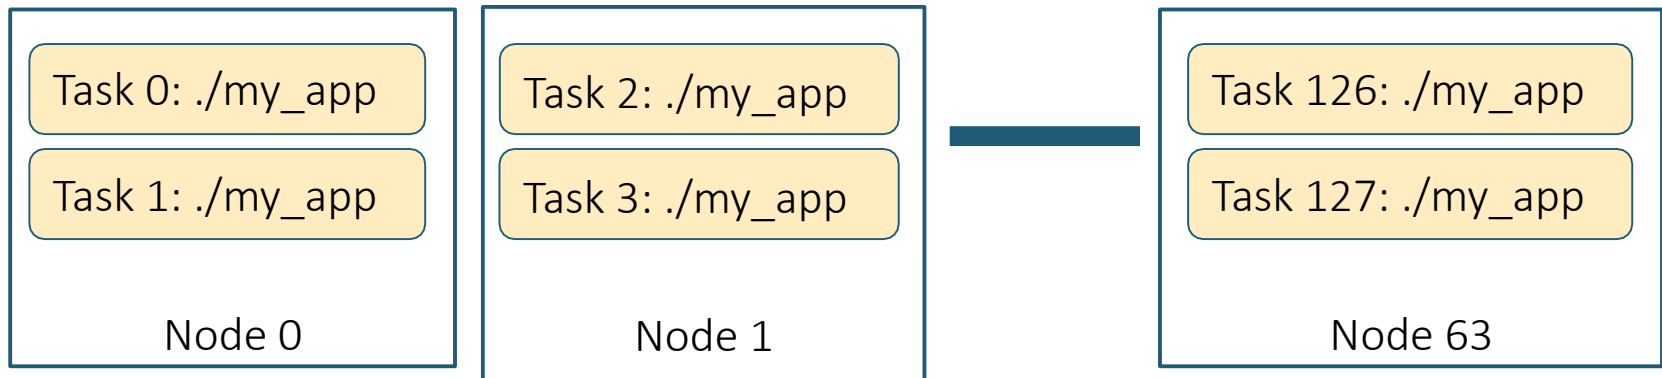
```
srun -N 32 ./my_app
```

```
srun -n 128 ./my_app
```

```
# request 64 nodes
```

```
# start ./my_app on 32 of them # (default: 1  
per node)
```

```
# start 128 instances of ./my_app,  
# across my 64 nodes (default is # to evenly  
distribute them in  
# block fashion)
```



One MPI rank generally corresponds to one SLURM Task

Requesting time

```
#SBATCH -t 30          # 30 minutes
#SBATCH -t 30:00       # 30 minutes
#SBATCH -t 1:00:00     # 1 hour
#SBATCH -t 1-0         # 1 day
#SBATCH -t 1-12        # 1.5 days
```

- Wallclock time, ie real elapsed time
- After this much time, SLURM can kill this job

Hands-on exercise: My first job

A SLURM job script has two sections:

1. Directives telling SLURM what you would like it to do with this job
2. The script itself - shell commands to run on the first compute node

```
elvis@nersc:~> vi myscript.q

#!/bin/bash -l

#SBATCH -t 00:30:00
#SBATCH -N 2
#SBATCH --license=SCRATCH

export RUNDIR=$SCRATCH/run-$SLURM_JOBID
mkdir -p $RUNDIR
cd $RUNDIR

srun -n 4 bash -c 'echo "Hello, world, from node $(hostname)"'

elvis@nersc:~> sbatch -C $CRAY_CPU_TARGET myscript.q
Submitted batch job 2774102
```

For how
long?

How many
nodes?

`$SCRATCH`
filesystem

Xeon nodes on
current cluster (set by
craype-
{haswell,ivybridge}
module)

Note: cannot use env
vars in directives - but
directives have

Hands-on exercise: My first job

A SLURM job script has two sections:

1. Directives telling SLURM what you would like it to do with this job
2. The script itself - shell commands to run on the first compute node

Make starting environment like my login environment

Run from \$SCRATCH

Start 4 tasks across my nodes

```
elvis@nersc:~> vi myscript.q
```

```
#!/bin/bash -l
```

```
#SBATCH -t 00:30:00
```

```
#SBATCH -N 2
```

```
#SBATCH --license=SCRATCH
```

```
export RUNDIR=$SCRATCH/run-$SLURM_JOBID
```

```
mkdir -p $RUNDIR
```

```
cd $RUNDIR
```

```
srn -n 4 bash -c 'echo "Hello, world, from node $(hostname)''
```

```
elvis@nersc:~> sbatch -C $CRAY_CPU_TARGET myscript.q
```

```
Submitted batch job 2774102
```

“sbatch”
submits a job
script

Running jobs – key points

- HPC work is via batch system
 - Dedicated subset of compute resources
 - Login nodes are shared resource for building code, editing scripts, etc. Use batch jobs for real work
- **Key commands:**
 - sbatch / salloc - submit a job
 - srun - start an (optionally MPI) application within a job
 - sqs - check the queue for my job status
- **Don't forget we have a reservation**
 - #SBATCH –reservation=csgrrain

<https://docs.nersc.gov/jobs/>

Where is my job?

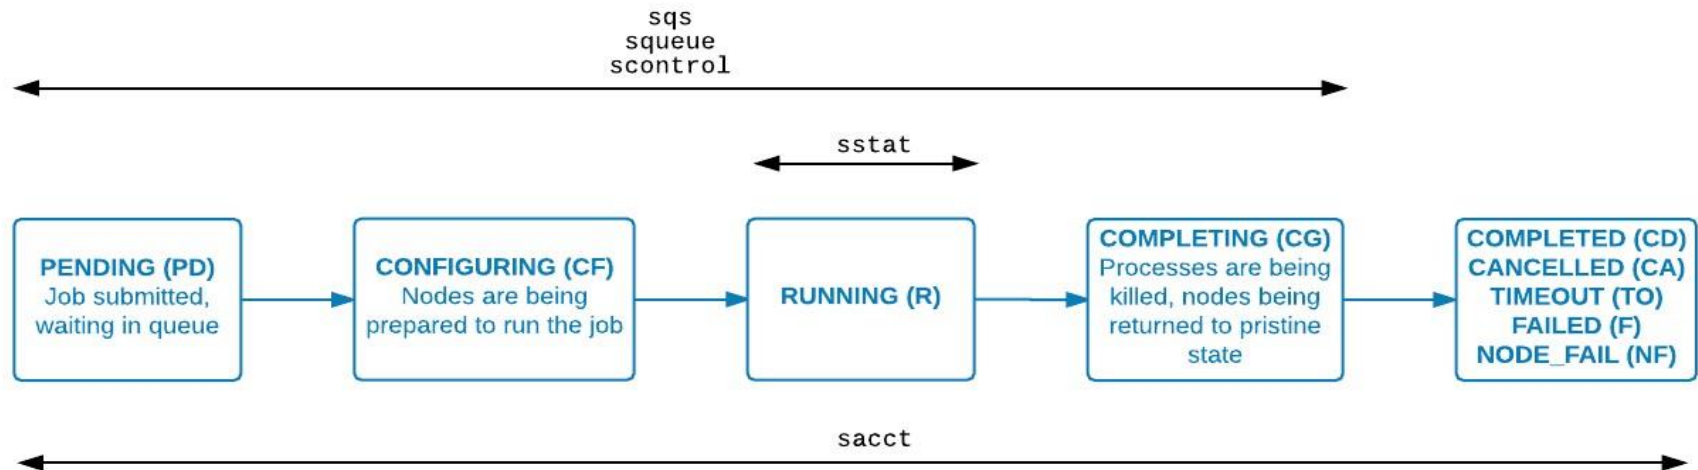
```
elvis@nersc:~> sqs
```

JOBID	ST	REASON	USER	NAME	NODES USED	REQUESTED ...
2774102	R	Prolog	elvis	myscript.q 2	0:00	30:00

...	SUBMIT	PARTITION	RANK_P	RANK_BF
	2016-11-18T11:24:20	debug	N/A	N/A

```
elvis@nersc:~> ls -lt total 11280
```

```
-rw-r----- 1 elvis elvis 132 Nov 18 11:24 slurm-2774102.out  
-rw-r----- 1 elvis elvis 208 Nov 18 11:24 myscript.q
```



Agenda

- Cori overview, logging in
- Run a simple job
- **Building and running applications on Cori**
 - Serial
 - Parallel (MPI)
 - Multithreaded (OpenMP)
- What affects performance?
 - Bottlenecks
 - Task placement and affinity
- Preparing for performance analysis

Cray compiler wrappers, ftn, cc and CC

- Building code optimally for Cori requires a complex set of compiler option and libraries
 - eg, static linking by default (important for performance at scale)
- Compiler wrappers ftn, cc and CC manage this complexity for you
 - Using environment variables set by the modules you have loaded
- Also provide MPI (so eg mpicc is not required)

Wait .. “environment modules” ?

- Software on Cori (and most HPC systems) is managed with “environment modules”
- Why?
 - Cori is a shared resource
 - Different people need different combinations of software, at different versions, with different dependencies (and for different jobs)
- Loading and unloading a module updates environment variables (eg \$PATH, \$LD_LIBRARY_PATH) to make a package available

Module Commands

module load <modulename>

- Add the module from your environment

module unload <modulename>

- Remove the module from your environment

module swap <module1> <module2>

- Unload one module and replace it with another

% module swap intel intel/16.0.3.210

(replace current default to a specific version)

module list

- See what modules you have loaded right now

module show <modulename>

- See what the module actually does

module display <modulename>

- Show module details

module help <modulename>

- Get more information about the software


Key modules for compiling

- PrgEnv-intel / PrgEnv-cray / PrgEnv-gnu
 - Which underlying compiler the wrappers should invoke
- craype-haswell / craype-mic-knl
 - Remember, login nodes are haswell, but we are building for KNL!

module swap craype-haswell

craype-mic-knl

- Wrappers manage cross-compiling



Important for today!

What do compiler wrappers link by default?

- Depending on the modules loaded, MPI, LAPACK/BLAS/ScaLAPACK libraries, and more

```
zz217@cori09:~/tests/dgemm> module list
```

```
Currently Loaded Modulefiles:
```

1) modules/3.2.10.5	7) udreg/2.3.2-4.6	13) job/1.5.5-3.58	19) craype-haswell
2) nsg/1.2.0	8) ugni/6.0.12-2.1	14) dvs/2.7_0.9.0-2.243	20) cray-shmem/7.4.4
3) intel/17.0.1.132	9) pmi/5.0.10-1.0000.11050.0.0.ari	15) alps/6.1.3-17.12	21) cray-mpich/7.4.4
4) craype-network-aries	10) dmapp/7.1.0-12.37	16) rca/1.0.0-8.1	22) altd/2.0
5) craype/2.5.7	11) gni-headers/5.0.7-3.1	17) atp/2.0.3	23) darshan/3.0.1.1
6) cray-libsci/16.09.1	12) xpmem/0.1-4.5	18) PrgEnv-intel/6.0.3	

```
zz217@cori09:~/tests/dgemm> ftn -v dgemmx.f -WL,-ydgemm_
```

```
...
```

```
ld /usr/lib64/gcc/x86_64-suse-linux/4.8/../../../../lib64/crt1.o /usr/lib64/gcc/x86_64-suse-linux/4.8/../../../../lib64/crti.o /usr/lib64/gcc/x86_64-suse-linux/4.8/crtbegin.o --build-id -static -m elf_x86_64 -L/opt/cray/pe/mpt/7.4.4/gni/sma/lib64 -L/opt/cray/pe/libsci/16.09.1/INTEL/15.0/x86_64/lib -L/opt/cray/dmapp/default/lib64 -L/opt/cray/pe/mpt/7.4.4/gni/mpich-intel/16.0/lib -L/opt/cray/dmapp/default/lib64 -L/opt/cray/pe/mpt/7.4.4/gni/mpich-intel/16.0/lib -L/usr/common/software/darshan/3.0.1.1/lib -L/opt/cray/rca/1.0.0-8.1/lib64 -L/opt/cray/alps/6.1.3-17.12/lib64 -L/opt/cray/xpmem/0.1-4.5/lib64 -L/opt/cray/dmapp/7.1.0-12.37/lib64 -L/opt/cray/pe/pmi/5.0.10-1.0000.11050.0.0.ari/lib64 -L/opt/cray/ugni/6.0.12-2.1/lib64 -L/opt/cray/udreg/2.3.2-4.6/lib64 -L/opt/cray/pe/atp/2.0.3/libApp -L/lib64 -L/opt/cray/wlm_detect/1.1.0-4.2/lib64 -o a.out /opt/intel/compilers_and_libraries_2017.1.132/linux/compiler/lib/intel64_lin/for_main.o -L/opt/intel/compilers_and_libraries_2017.1.132/linux/compiler/lib/intel64 -L/opt/intel/compilers_and_libraries_2017.1.132/linux/mkl/lib/intel64 -L/opt/intel/compilers_and_libraries_2017.1.132/linux/compiler/lib/intel64_lin -L/usr/lib64/gcc/x86_64-suse-linux/4.8/ -L/usr/lib64/gcc/x86_64-suse-linux/4.8/../../../../lib64 -L/usr/lib64/gcc/x86_64-suse-linux/4.8/../../../../lib64/ -L/lib/./lib64 -L/lib/./lib64/ -L/usr/lib/./lib64 -L/usr/lib/./lib64/ -L/opt/intel/compilers_and_libraries_2017.1.132/linux/compiler/lib/intel64/ -L/opt/intel/compilers_and_libraries_2017.1.132/linux/mkl/lib/intel64/ -L/usr/lib64/gcc/x86_64-suse-linux/4.8/../../../../x86_64-suse-linux/lib/ -L/usr/lib64/gcc/x86_64-suse-linux/4.8/../../../../lib64 -L/lib/ -L/lib/ -L/usr/lib64 -L/usr/lib /tmp/iftortsklewd.o -ydgemm_ @/usr/common/software/darshan/3.0.1.1/share/ld-opts/darshan-base-ld-opts -lfmpich -lmpichcxx --start-group -ldarshan -ldarshan-stubs --end-group -lz --no-as-needed -lAtpSigHandler -lAtpSigHCommData --undefined=ATP_Data_Globals --undefined=__atpHandlerInstall -lpthread -lmpichf90_intel -lrt -lugni -lpmi -L/opt/intel/compilers_and_libraries_2017.1.132/linux/compiler/lib/intel64_lin -limf -lm -lpthread -ldl -lsma -lpmi -lsma -lpmi -ldmapp -lpthread -lsci_intel_mpi -lsci_intel -L/opt/intel/compilers_and_libraries_2017.1.132/linux/compiler/lib/intel64_lin -limf -lm -ldl -lmpich_intel -lrt -lugni -lpthread -lpmi -L/opt/intel/compilers_and_libraries_2017.1.132/linux/compiler/lib/intel64_lin -limf -lm -ldl -lpmi -lpthread -lalpslli -lpthread -lwlm_detect -lalpsutil -lpthread -lrca -lxpmem -lugni -lpthread -ludreg -lsci_intel -L/opt/intel/compilers_and_libraries_2017.1.132/linux/compiler/lib/intel64_lin -limf -lm -ldl --as-needed -limf --no-as-needed --as-needed -lm --no-as-needed --as-needed -lpthread --no-as-needed -lifort -lifcore -limf -lsvml -lm -lipgo -lirc -lsvml -lc -lgcc -lgcc_eh -lirc_s -ldl -lc /usr/lib64/gcc/x86_64-suse-linux/4.8/crtend.o /usr/lib64/gcc/x86_64-suse-linux/4.8/../../../../lib64/crtn.o /tmp/iftortsklewd.o: reference to dgemm_ /opt/cray/pe/libsci/16.09.1/INTEL/15.0/x86_64/lib/libsci_intel.a(dgemm.o): definition of dgemm_
```

Compiling code

- Very similar to regular Linux, but using CC / cc / ftn
- Do this bit once:
module swap craype-haswell craype-mic-knl
- Then:
cc -c hack-a-kernel.c
cc -o hack-a-kernel.ex hack-a-kernel.o
- Note that the module looks apear CPU target!

Compiling parallel code

- Compiler wrappers give you MPI “for free”

```
CC -c hello-mpi.c++
```

```
CC -o hello-mpi.ex          hello-mpi.o
```

- (Cray MPICH – optimized for Aries HSN)

- OpenMP: with PrgEnv-intel (NERSC default):

```
cc -qopenmp -c hello-omp.c
```

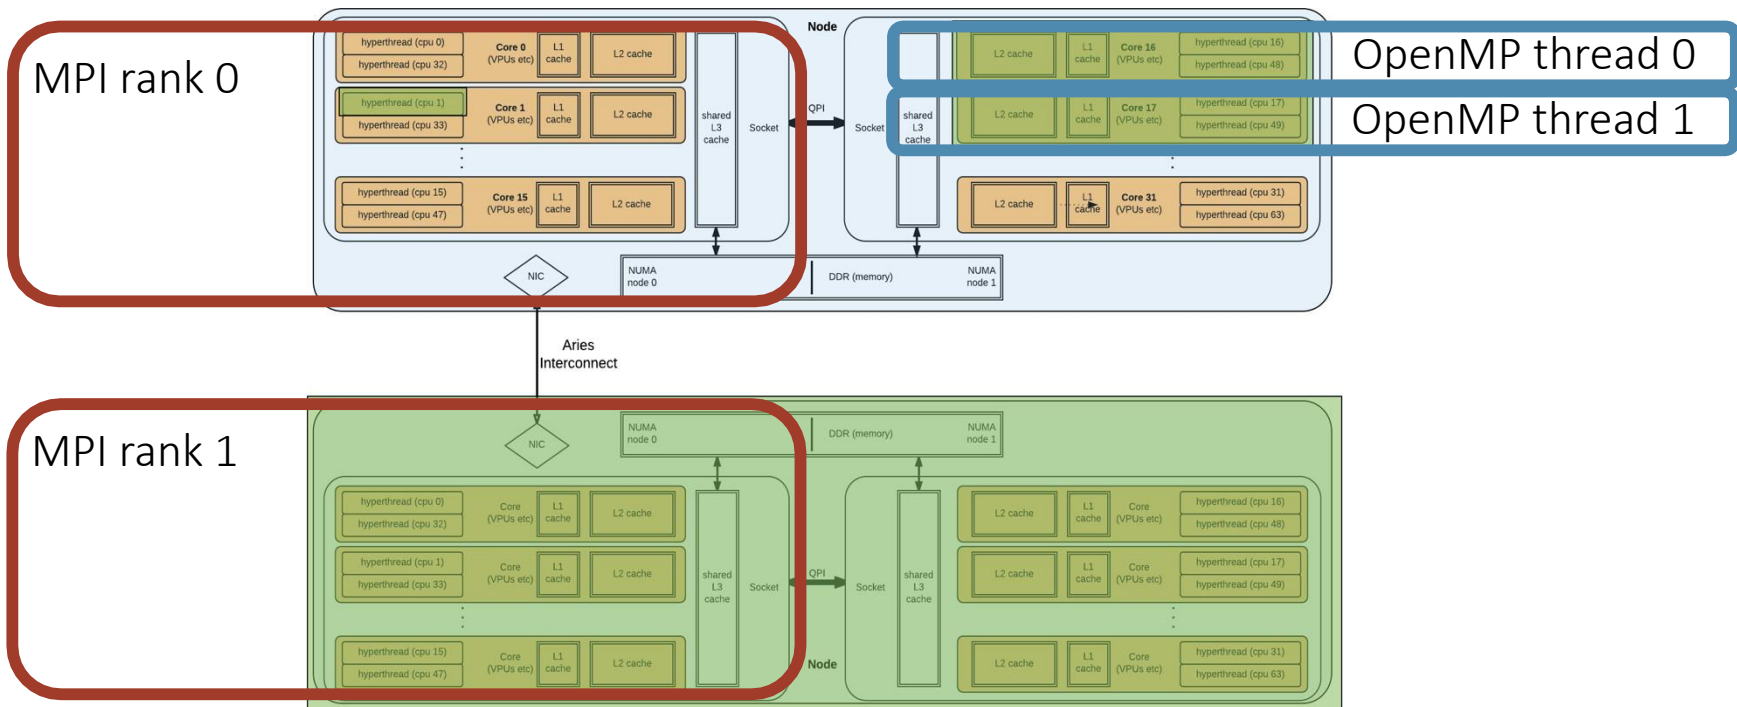
```
cc -qopenmp -o hello-omp.ex hello-omp.o
```

MPI vs OpenMP

- **MPI provides explicit communication between separate processes**
 - Optionally on separate nodes – ie packets over a network
 - Most parallel development in last 2 decades has used this approach
- **OpenMP provides work-sharing and synchronization between threads in a single process**
 - Threads share the same memory image
 - To make the most of a KNL node, most applications will need to use OpenMP

MPI vs OpenMP


- An MPI application can have processes on more than one node
- An OpenMP application exists entirely within 1 node



Why do we suddenly need OpenMP?

- Multi-level parallelism

- At very large scale, the overheads of MPI (or any parallel approach) become excessively costly
- Combining (nesting) parallel approaches allows us to operate each at lower scale
 - Sweet spot for best overall efficiency
- MPI -> OpenMP -> Vectorization



Why do we suddenly need OpenMP?

- Memory-per-core is trending downwards
 - Cori Haswell: 128GB for 32 cores
 - Cori KNL: 96GB for 68 cores (16GB MCDRAM for 68 cores)
 - Parallelism within same memory footprint is necessary

Agenda

- Cori overview, logging in
- Run a simple job
- Building and running applications on Cori
 - Serial
 - Parallel (MPI)
 - Multithreaded (OpenMP)
- **What affects performance?**
 - Bottlenecks
 - Task placement and affinity
- Preparing for performance analysis

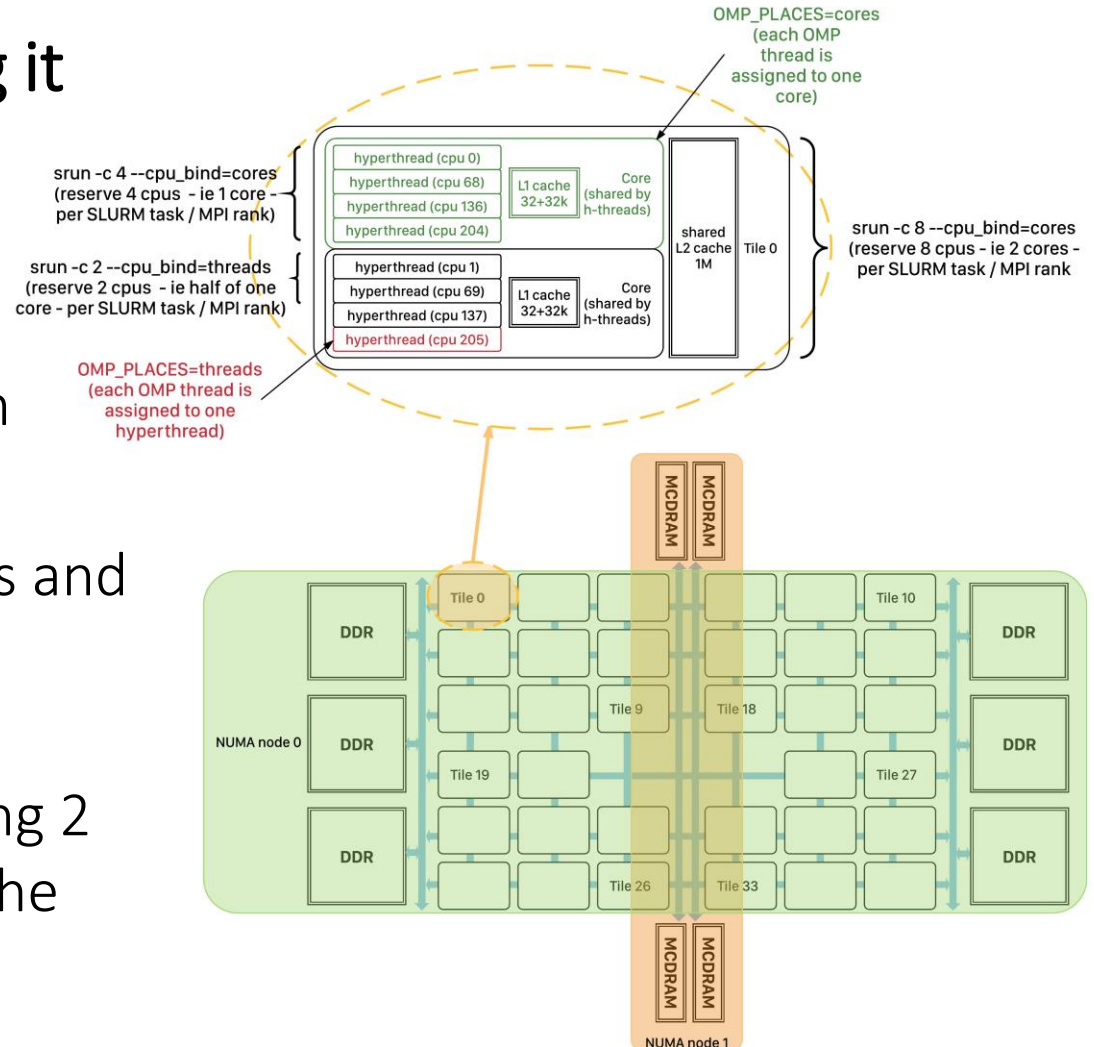
Performance bottlenecks

- Any point at which some component of the system or application is stalled, waiting on some other component, is a bottleneck
- Eg waiting for MPI call to complete
 - Load imbalance or communication overhead is a bottleneck
- BUT for today we are interested in KNL specifics
 - Bottlenecks within the node

Bottlenecks within the node – Affinity issues

- Firstly: Am I running it right?

- Cori KNL has 68 cores per node, arranged on a mesh of 34 tiles
- Each tile has 2 cores and a shared L2 cache
- Each core has 4 hyperthreads sharing 2 VPU and an L1 cache



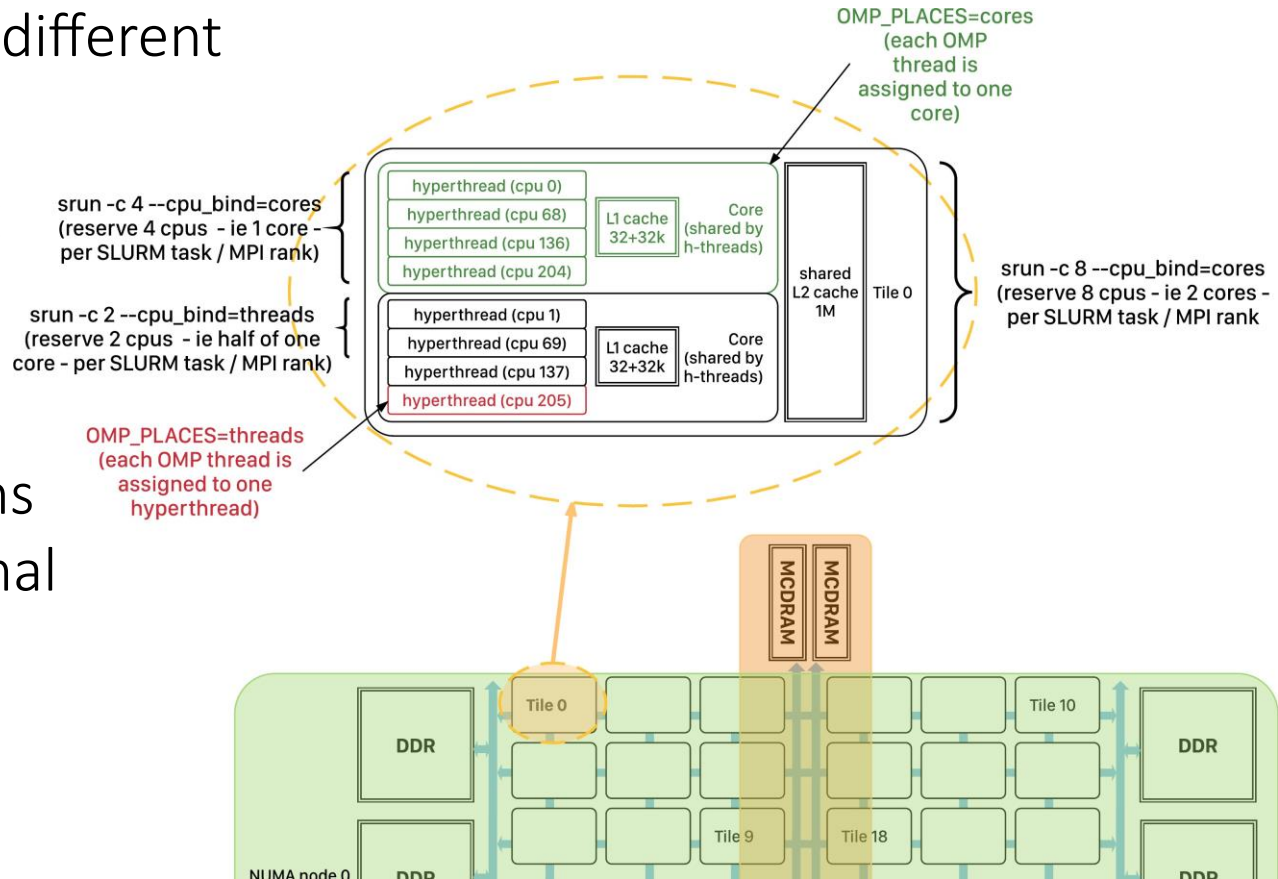
Bottlenecks within the node – Affinity issues

- Where are my threads?

- Is each using a different core at least?

- Linux does not always choose best placement

- Use srun options to ensure optimal thread/process placement



Process (task) affinity

Solution: use `--cpu_bind`:


`srun -n 64 -c 4 --cpu_bind=verbose,cores`

`./my_exec`


`srun -n 128 -c 2 --cpu_bind=verbose,threads`

`./my_exec`

- Controls what a task (MPI rank) is bound to
 - If no more than 1 MPI rank per core: `--cpu_bind=cores`




- If more than 1 MPI rank per core: `--cpu_bind=threads`



Thread affinity (OpenMP)

```
export OMP_NUM_THREADS=2 export  
OMP_PROC_BIND=spread export      # or close  
OMP_PLACES=cores                  # or threads, or sockets  
srun -n      32 -c 8 --cpu_bind=verbose,cores      ./my_exec
```



...If using hyperthreads, use OMP_PLACES=threads

Memory affinity

Linux default behavior is to allocate to closest NUMA-node, if possible

Not always optimal:

- KNL nodes: DDR is “closer” than MCDRAM

```
#SBATCH -C knl,quad,flat export
```

```
OMP_NUM_THREADS=4
```

```
srun -n16 -c16 --cpu_bind=cores --mem_bind=map_mem:1 ./a.out
```

- NUMA node 1 is MCDRAM in quad,flat mode
- “Mandatory” mapping: if using >16GB, malloc will fail

NOTE: today’s reservation is for “cache-mode” nodes, so MCDRAM is invisible

Bottlenecks within the node – Affinity issues

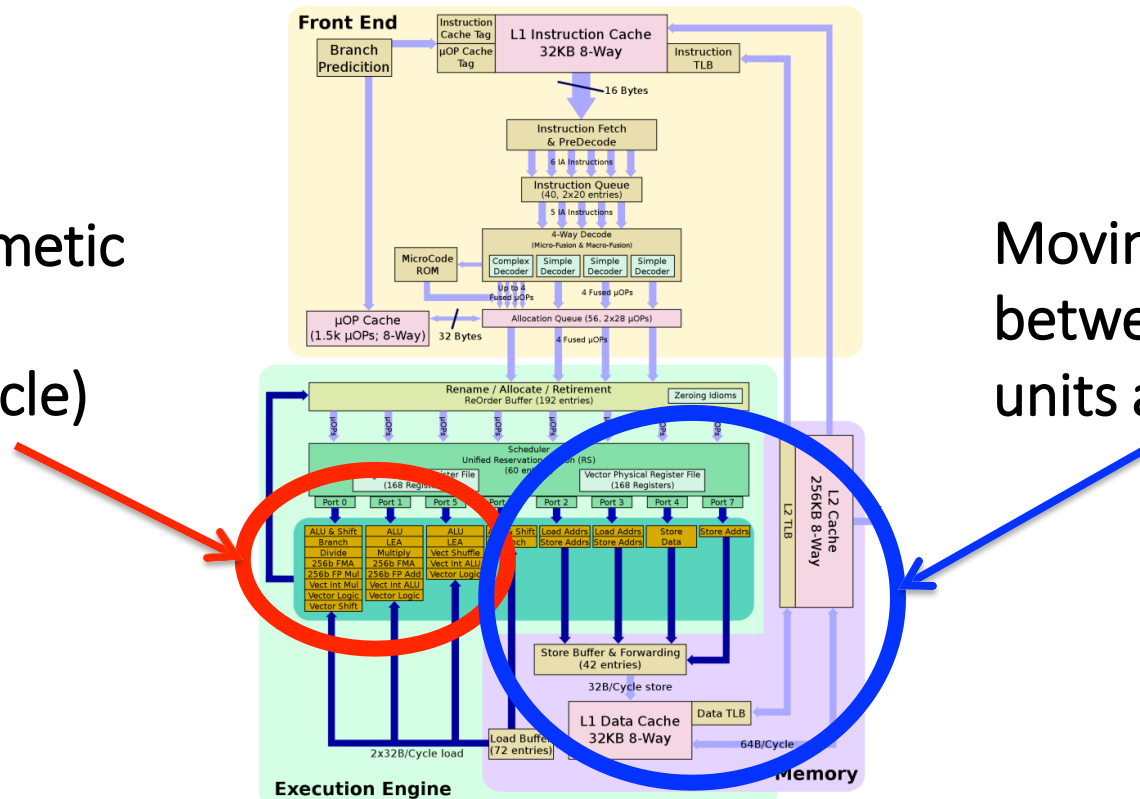
OMP_PROC_BIND=true OMP_PLACES=cores (threads) srun -c 4 --cpu_bind=cores ...

Bottlenecks within the node

- Performance tends to be dominated by:

Doing arithmetic operations
(# FLOPS/cycle)

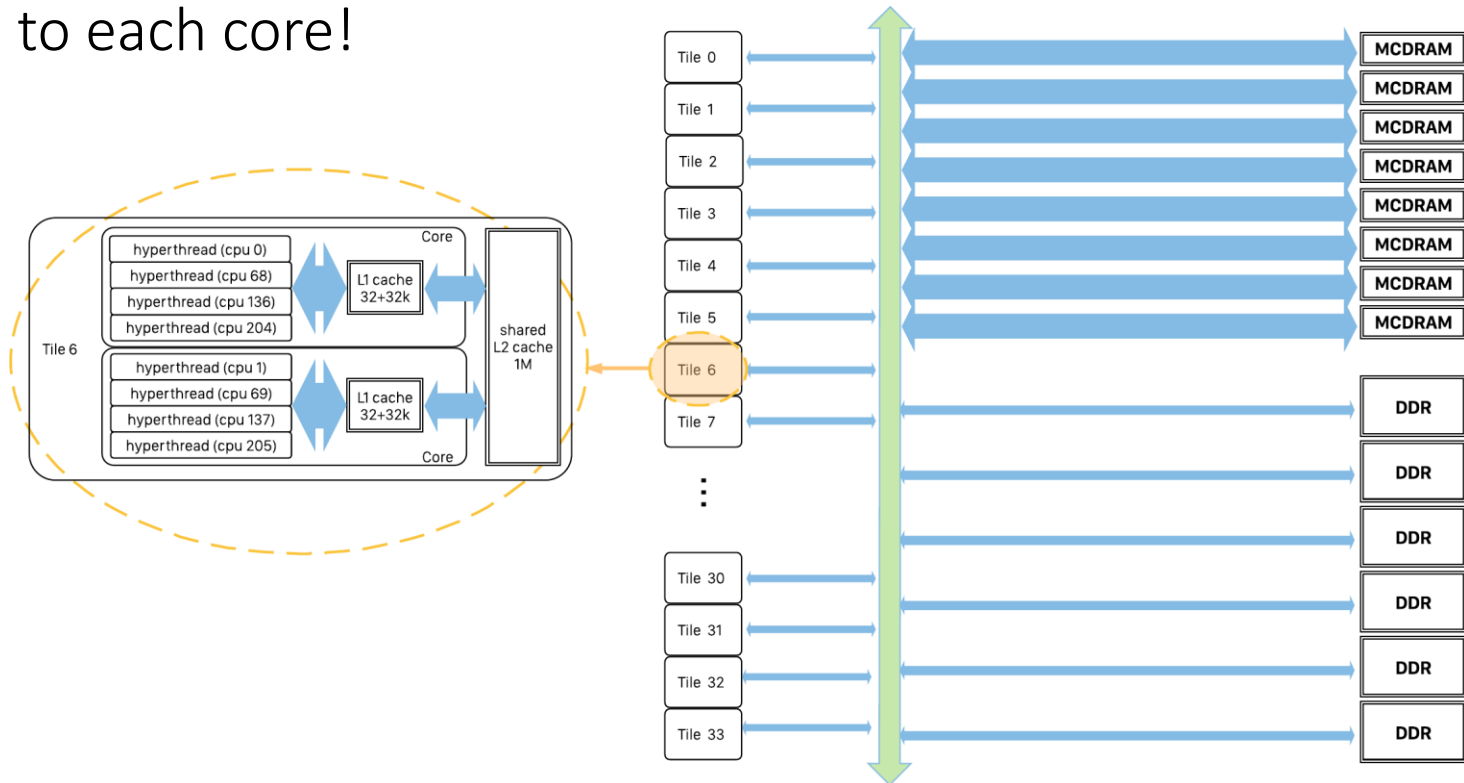
Moving data between arithmetic units and memory



- Bottlenecks are usually due to one of these parts waiting on the other

Bottlenecks within the node

- Bandwidth!
 - MCDRAM is very high bandwidth (~ 450 GB/s)
 - .. But aggregate bandwidth is not the same as bandwidth to each core!




Agenda

- Cori overview, logging in
- Run a simple job
- Building and running applications on Cori
 - Serial
 - Parallel (MPI)
 - Multithreaded (OpenMP)
- What affects performance?
 - Bottlenecks
 - Task placement and affinity
- **Preparing for performance analysis**

Profiling prerequisite

- A small, short, but representative test case for your application
 - Profiling tends to be costly
 - Runtime overhead
 - Size of data collected
 - **BUT:** must cover same paths through code as “real” example (or at least, the differences must be understood)
 - What could go wrong with this test case?



...But my application *just is* big!

- Start with a low-overhead profiling method
 - Eg sampling-based (gprof, TAU, CrayPat, ...)
 - Identify hotspots
- Only profile part of a run
 - Some tools (eg Vtune) allow you to start the run “paused” and “resume collection” via an API call
- Only profile 1 MPI rank
 - Via srun options, eg run Vtune on one rank, not others (beyond today’s scope)
- Luckily, our hack-a-kernel is already small!

Performance Analysis Gotchas

- Many tools require dynamically-linked executable
 - Including Vtune
 - NERSC/Cray: applications are statically linked by default, must use “-dynamic” at compile time
- “Uncore” (eg memory related) performance counters usually require special permissions (eg kernel module)
 - NERSC/Slurm supports this with `#SBATCH --perf=vtune` directive

Summary and Recap

- Now you can:
 - Log in
 - Build an application
 - Run an application, and be aware of how and where it is running
 - Prepare an application for performance analysis with Vtune
- Q & A?



NERSC